Time Triggered Offline Scheduler for Data Dependent Real-Time Tasks Accounting for Preemption and Scheduler Costs

Walid Talaboulma
INRIA Paris-Rocquencourt
Domaine de Voluceau BP 105
78153 Le Chesnay Cedex - France
walid.talaboulma@inria.fr

Falou Ndoye
INRIA Paris-Rocquencourt
Domaine de Voluceau BP 105
78153 Le Chesnay Cedex - France
falou.ndoye@inria.fr

Yves Sorel
INRIA Paris-Rocquencourt
Domaine de Voluceau BP 105
78153 Le Chesnay Cedex - France
yves.sorel@inria.fr

Abstract—**Time critical embedded systems usually consist of a set of periodic data dependent real-time tasks which exchange data. Although non-preemptive real-time scheduling is safer than preemptive real-time scheduling in a time critical context, preemptive real-time scheduling provides a better success ratio. However, the preemption has a cost that it is risky to take into account imprecisely. In this paper we propose a schedulability analysis for data dependent periodic tasks which precisely accounts for preemption and scheduler costs. It results in a scheduling table that is exploited in a time triggered offline scheduler. We show that this scheduler, implemented on an ARM Cortex-M4 bare metal processor, is able to schedule correctly set of tasks that miss their deadline when preemption and scheduler costs are neglected. Therefore, such a scheduler is perfectly suited for time critical embedded systems.**

Keywords-time critical embedded systems, real-time scheduling, data dependent tasks, preemption cost, scheduler cost, offline schedulability analysis, time triggered scheduler, ARM Cortex-M4.

I. INTRODUCTION

We address time critical embedded systems, i.e. systems for which time constraints must necessarily be satisfied in order to avoid catastrophic consequences. Such systems, in most cases, consist of a set of data dependent periodic tasks resulting from a functional specification, usually achieved with tools such as Simulink [1], Scade [2], etc., based on block diagrams. The functional specification describes the functions that must be executed, as well as their dependences carrying the data produced and consumed by the functions. Such dependences involve a precedence relation on the execution of every producer function relatively to one or several consumer functions, and lead to sharing the transfered data. Data dependent functions associated with temporal characteristics become data dependent real-time tasks. Some of these characteristics like first releases, periods and deadlines are not related to the processor that will execute the real-time tasks, whereas the values of WCET (Worst Case Execution Time) depend on the processor. Usual schedulability analyses of periodic data dependent tasks are based on the WCET, and thus, such analyses require that the designer determines accurately the WCET

which is strongly related to the internal architecture of the processor. More this architecture is complex more it is difficult to determine the WCET.

Although non-preemptive real-time scheduling is safer than preemptive real-time scheduling in a time critical context, preemptive real-time scheduling provides a better success ratio. However, the preemption has a cost that it is risky to take into account imprecisely. The usual way to account for this cost, consists in adding to the WCET of every task an extra cost which is a percentage of this WCET. This approach may lead to miss some deadlines during the runtime execution of the tasks even though the schedulability conditions have been satisfied or, in the best case, resources could be wasted when this percentage is chosen to high.

Therefore, on the one hand in order to guarantee the real-time schedulability of a set of periodic data dependent tasks specifying a critical embedded systems and on the other hand to minimize its needed resources, we propose a schedulability analysis which takes into account precisely the cost of preemption by counting them along a study interval. This analysis produces a scheduling table that is exploited in a time triggered offline scheduler. We give the principles of this scheduler and implement it on an ARM Cortex-M4 bare metal processor. Then, whe show that this implementation is able to schedule correctly a set of tasks that are not correctly scheduled when preemption and scheduler costs are neglected.

The remainder of the paper is organized as follows. Section II presents the related work on periodic data dependent tasks and on the preemption cost. Section III presents the schedulability analysis. Section IV gives the principles of the time triggered offline scheduler. Section V presents a performance evaluation of the proposed scheduler on an ARM Cortex-M 4 bare metal processor. Finally, Section VI concludes and gives some directions for future work.

II. RELATED WORK

We assume that the processor where the tasks will be executed, has neither cache nor complex pipeline or specific internal architecture features. The previous assumptions are

usually made in time critical embedded systems where determinism is a key issue. In this case, the preemption cost corresponds to the duration necessary to save the context of the preempted task and the duration necessary to restore this context when the preempted task will be selected again to resume its execution. Due to its cost, a preemption increases the response time of the preempted task that may cause another preemption, and so on. The cost of the preemption is usually approximated in the WCET as assumed, explicitly, by Liu and Layland in their pioneering article [3]. That is, some percentage of the WCET which corresponds to the longest path in the sequential program associated to a task, is added to its actual value. When the preemption cost is neglected, meaning this percentage is close to zero, although a set of tasks verifies the schedulability condition associated to the scheduling algorithm - which will be implemented in the real-time scheduler -, some deadline misses may occur. In order to tackle this problem, a first solution consists in determining the maximum number of preemptions as proposed in [4] or in determining the number of preemptions but without accounting for the cost of each preemption that can cause other preemptions, increasing the global cost in [5]. Other solutions aim at controlling the number of preemptions, like presented in [6]. It is worth noting that the preemption cost is not the only cost that must be precisely accounted when dealing with time critical systems. Indeed, the scheduler cost itself must be also precisely accounted. Taking into account the maximum number of preemptions and the scheduler cost, lead to increase the WCET up to 50%, for example in the most critical applications of the avionic industry. This pessimism decreases the schedulability ratio and increases the amount of necessary resources. On the other hand, a solution that determines the exact number of preemptions while accounting for the cost of each preemption is proposed in [7]. Unfortunately, this solution assumes that the scheduling algorithm is based on fixed priorities.

Periodic data dependent tasks mean there are precedence constraints between tasks [8] such that a producer task is executed before the corresponding consumer task, and the consumer task must receive the data produced by the producer task. There are two approaches to dealing with precedence constraints. The first one is based on semaphores [9]. A semaphore is allocated to each precedence, and the consumer task must wait for the producer task to release the semaphore before it can start its execution. The second approach is based on the modification of the priorities and the release times of the task [8], [10]. Actually, when dependent tasks have different periods the problem is much more complex than when they have the same period [11]. If the producer task τ_p has a period smaller than the consumer task τ_c, then the latter has to consume in the worst case $n = \lfloor \frac{\tau_c}{\tau_p} \rfloor$ data produced by the producer task. This worst case approach was chosen in [12]. Usually, only the last data

is consumed since it is considered to be the "freshest" one, like in [11]. Conversely, when the producer task has a period greater than the consumer task, the latter has to consume at worst n times the same data produced by the producer task. Thus, it is sufficient that the consumer task consumes only one of these data. When tasks are data dependent they have to share some buffer containing the data that may involve priority inversions. For example, a lower priority producer task can block the execution of a consumer task that want to read it while it has a higher priority. In order to avoid this situation the well known priority inheritance protocol that gives to a task the highest priority of all the tasks which share a data, was proposed in [13]. This protocol holds only for static priority scheduling algorithms. It was extended in [13] by giving an additional priority to every shared data equal to the highest priority of the tasks that share this data. This priority ceiling protocol minimizes the blocking time and prevents deadlocks that may occur when several tasks are mutually waiting for a shared data used by other tasks. For dynamic priority scheduling algorithms, the stack resource policy was proposed in [14].

In order to take into account precisely preemption and scheduler costs, an offline schedulability analysis that considers the cost of each preemption, is proposed in [15]. Moreover, this analysis allows changes in the priorities of the tasks that are necessary for dependent tasks which involve priority inversions. In this paper after summarizing the principles of this schedulability analysis, we show how it is exploited to implement an offline scheduler executed at runtime that is deterministic and thus perfectly suited for time critical embedded systems.

III. SCHEDULABILITY ANALYSIS

The schedulability analysis is based on a schedulability interval. This is a finite time interval such that the schedule on this interval can be repeated infinitely. We use the minimal schedulability interval for a set of n periodic data dependent tasks proposed in [16]. I_n denotes the schedulability interval, given by:

$$I_n = [r_{min}, t_c + H_n] \qquad (1)$$

where t_c denotes the time from which the schedule repeats indefinitely. t_c is computed iteratively by an algorithm given in the article. Since t_c is smaller or equal to $r_{max} + H_n$, we will use thereafter the interval given by the Equation 1 with $tc = r_{max} + H_n$, r_{min} and r_{max} are respectively the minimum and the maximum of the first release times r_i^1 of the tasks τ_i which is released at times r_i^k, i.e. every instance (job) $k = 1..\infty$ of the task.

Γ_n denotes the set of periodic dependent tasks. The schedulability analysis of Γ_n is achieved on the schedulability interval I_n according to a given fixed or dynamic priority scheduling algorithm, for example Rate Monotonic

(RM) or Earliest Deadline First (EDF) [3], to cite only the most famous ones. Moreover, the release times and the deadlines of every task are modified such that a task τ_j can be executed if and only if each of its predecessors τ_i produces $k_{ij} = \lceil \frac{T_j}{T_i} \rceil$ data, and τ_j does not produce no more than $k_{jk} = \lceil \frac{T_k}{T_j} \rceil$ data for each of its successors τ_k. These conditions guarantee that all the data produced are consumed when two dependent tasks have equal or different periods and prevent deadlocks between tasks. However, these conditions do not prevent priority inversions due to the data shared by the dependent tasks. This is the reason why we use, in the given scheduling algorithm, the priority inheritance protocol [13] which minimizes the duration of priority inversions.

The schedulability analysis performs a simulation of an offline scheduler that is called only at release and completion times of every task. For every of theses calls, denoted t, the schedulability analysis selects among the ready tasks, with the function denoted $\phi : I(t) \rightarrow \Gamma_r(t)$, the task to execute denoted τ_i. Then, it computes the remaining execution time of τ_i denoted $c_i : I(t) \rightarrow \mathbb{N}$ and the relative deadline of τ_i denoted $d_i : I(t) \rightarrow \mathbb{N}$. These three functions are used to test the schedulability of the task τ_i. Finally, it determines the next scheduler call. At the end of the schedulability interval I_n, if $\forall \tau_i \in \Gamma_n$, τ_i is schedulable then Γ_n is schedulable and a scheduling table is produced, else Γ_n is not schedulable.

A. Task Selection $\phi(t)$

When the scheduler is called at t, the task to be selected must belong to the ready set of tasks denoted $\Gamma_r(t)$. A task τ_i is ready at t if an only if: its first release time occurs before, or at t, and it received all the data produced by its predecessors, and all its successors consumed all the data it produced. The selected task, denoted $\phi(t)$, is the task with the highest priority in $\Gamma_r(t)$ according to the given scheduling algorithm and the priority inheritance protocol.

B. Remaining Execution Time $c_i(t)$

$c_i(t)$ is the number of time units that τ_i must still execute at t to complete its execution. If τ_i is preempted at t, the cost of one preemption is added to the remaining execution time $c_i(t)$ of τ_i.

At every release or completion time of a task τ_i its remaining execution time $c_i(t)$ is given by:

$$
c_i(t) = \begin{cases}
C_i & \text{if } (\frac{t - r_i^1}{T_i}) \in \mathbb{N} \text{ else} \\[2mm]
c_i(r^-(t)) & \text{if } (\phi(r^-(t)) \neq \tau_i) \text{ else} \\[2mm]
c_i(r^-(t)) - (t - r^-(t)) & \text{if } (\phi(t) = \tau_i) \vee \\
& ((\phi(t) \neq \tau_i) \wedge \\
& (r^-(t) + c_i(r^-(t)) = t)) \text{ else} \\[2mm]
c_i(r^-(t)) - (t - r^-(t)) + \alpha
\end{cases}
$$

where C_i denotes the WCET of τ_i, α the cost of one preemption, and $r^-(t)$ the previous scheduler call. It is important to note that the WCET is considered here without any approximation of the preemption cost since this cost is precisely taken into account with α. However, the WCET includes the precise cost for storing the context when a task is released while preempting another task. In addition, it includes the precise cost of the scheduler, as mentioned in the section II, which is very simple in our case and can be deterministically and precisely determined. This cost consists in reading, in the scheduling table, the next task to execute when it is released, or when it is resumed if this task were preempted. This is the cost of the interruption routine given in section IV-B where we present the runtime scheduler algorithm 1.

In this computation there are four cases: 1) this is τ_i which is released at t and thus $c_i(t) = C_i$, 2) during the previous scheduler call the selected task was different from τ_i and thus the remaining execution time of τ_i does not change $c_i(t) = c_i(r^-(t))$, 3) during the previous scheduler call the selected task was τ_i and it is not preempted at t, meaning that τ_i is still the selected task at t or, that τ_i completes its execution, thus $c_i(t) = c_i(r^-(t)) - (t - r^-(t))$. That is, the time elapsing between t and the previous value of t corresponding to the execution time of τ_i, is subtracted to the previous value of c_i, 4) during the previous scheduler call the selected task was τ_i and τ_i is preempted at t, it follows that the cost of one preemption α is added to $c_i(r^-(t)) - (t - r^-(t))$.

Of course, this approach accounts for the cost of each preemption that can lead to produce other preemptions.

The figure 1 shows an example with two periodic tasks $\tau_1(2, 2, 6, 6)$ and $\tau_2(0, 3, 8, 8)$ where the timing characteristics between brackets correspond respectively to the first release time, the WCET, the deadline, and the period. We assume that the scheduling algorithm is RM. In this example, the offline scheduler is called at t equal 0, 2, 4, etc., corresponding to the release and completion times of both tasks τ_1 and τ_2. At $t = 0$, τ_2 is released, thus $c_2(0) = 3$. Then, at $t = 2$, τ_1 is released, thus $c_1(2) = 2$ and τ_2 is preempted by τ_1. A time unit (in black) is added to take into account for the cost for restoring the context of τ_2 while the cost for storing the context of τ_2 is assumed to be included in the WCET of τ_1, thus $c_2(2) = (3 - 2 + 0) + 1 = 2$. For the sake of simplicity, in this example we chose one time unit for the cost for restoring the context, but actually it is widely smaller than the WCET. For the same reason we do not show the cost for storing the context in the WCET of the preempting task. See section **??** to have an idea of the realistic values of these costs. At $t = 4$, τ_1 completes its execution, thus $c_1(4) = 2 - 4 + 2 = 0$ and τ_2 resumes. Since during the previous scheduler call τ_1 was selected, which is different of τ_2, thus $c_2(4) = c_2(2) = 2$, and so one for the other scheduler calls.
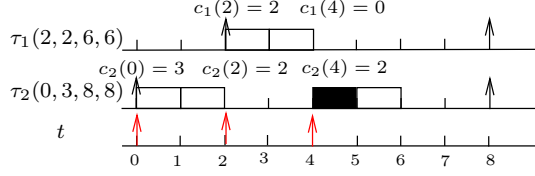
Figure 1. Remaining execution time accounting for preemption cost

C. Relative Deadline $d_i(t)$

At every release of τ_i, $d_i(t) = D_i$, then $(t - r^-(t))$ is substracted to $d_i(t)$ every time the scheduler is called. In order to do not miss its deadline τ_i must complete its execution before date $t + d_i(t)$.

$d_i(t)$ is given by:

$$
d_i(t) = \begin{cases}
D_i & \text{if } (\frac{t-r_i^1}{T_i}) \in \mathbb{N} \quad \text{else} \\
d_i(r^-(t)) - (t - r^-(t)) & \text{if } r^-(t) + d_i(r^-(t)) > t \\
& \text{else} \\
0 &
\end{cases}
$$

In this computation there are three cases: 1) the task τ_i is released at t, 2) the previous scheduler call added to the previous relative deadline is greater than the present scheduler call, thus the time elapsing between t and the previous value of t, is substracted to the previous value of d_i, 3) the task τ_i completes or has already completed.

The figure 2 shows an example with one task $\tau_i(0,3,8,8)$. At $t = 0$ the task τ_i is released, for the first time, thus $d_i(0) = 8$. At $t = 2$, $r^-(2) - d_i(r^-(2)) = 0 + 8 > 2$, thus its relative deadline $d_i(2) = 8 - 2 + 0 = 6$ since its previous value was $d_i(0) = 8$.
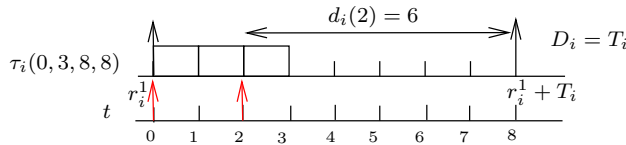

Figure 2. Relative deadline

D. Schedulability Condition at t

The schedulability condition guarantees that the set of tasks is schedulable at every scheduler call t. According to the theorem given in [15], a task $\tau_i \in \Gamma_n$ is schedulable at t if and only if:

$$
\begin{aligned}
& (c_i(t) \leq d_i(t)) \wedge ((t \leq r_i^1) \vee (c_i(r^-(t)) = 0) \vee \\
& (\phi(r^-(t)) = \tau_i) \vee ((t - r_i^1) mod T_i \neq 0))
\end{aligned} \tag{2}
$$

where T_i is the period of the task τ_i.

This condition means that, at t, the remaining execution time of τ_i is less than or equal to its deadline and one of the following cases occurs: τ_i is not still released, or τ_i completes its execution, or τ_i was the selected task at the previous scheduler call, or τ_i does not begin a new instance without completing its execution in its previous instance.

Therefore, the set of tasks Γ_n is schedulable at t if and only if $\forall \tau_i \in \Gamma_n$, τ_i is schdedulable at t.

Moreover, this schedulability condition is sustainable according to the WCET. That is, even if some tasks have execution times smaller than their WCET then the set of tasks remains schedulable. This is an important property when this approach is actually implemented on a bare metal processor, as it will be presented in the next sections.

E. Next Scheduler Call $r^+(t)$

The next $r^+(t)$ scheduler call corresponds to a release or a completion of a task belonging to Γ_n.

$r^+(t)$ is given by:

$$
r^+(t) = \begin{cases}
t + c_j(t) & \text{if } ((t + c_j(t)) < r(t)) \wedge (\phi(t) = \tau_j) \text{ else} \\
r(t) &
\end{cases}
$$

where $r(t)$ denotes the next release time of a task that belongs to the set $F = \{t \in I_n / \exists (\tau_i, k) \in (\Gamma_n, \mathbb{N}), t = r_i^1 + kT_i\}$ containing the release times in I_n of the set of tasks Γ_n. $r(t)$ is the successor element t in F.

In this computation there are two cases: 1) the selected task at t is τ_j and its remaining execution time $c_j(t)$ added to t is less than the next release time of a task, i.e. the next scheduler call corresponds to the completion time of τ_j, 2) the next scheduler call corresponds to the next release time of a task.

F. Schedulability Analysis Algorithm

The schedulability analysis of the set of task Γ_n is performed by the algorithm 1. It moves iteratively through the elements t of the schedulatibilty interval I_n which contains only release and completion times of the tasks corresponding to the scheduler calls. For every time t it selects the task to execute and verifies if there is a non schedulable task using the schedulability condition 2. As soon as a task is not schedulable the set of task Γ_n is not schedulable. Otherwise if all the tasks verify the condition 2 the set of task Γ_n is schedulable. In this case a scheduling table is produced containing, for every scheduler call, the task to execute and a status indicating if the task is released or resumed.

As the schedulability analysis is performed offline, according to a fixed or dynamic priority scheduling algorithm, fairly complex task sets can be handled. Should a feasible solution not be found, retries are possible, e.g., by changing the parameters of the scheduling algorithm or the timing characteristics of the task set.

Algorithm 1 Schedulability analysis

1: $t \leftarrow r_{min}$
2: $G \leftarrow F$
3: $schedulable \leftarrow true$
4: **while** $(t < (tc + H_n)) \wedge (schedulable = true)$ **do**
5:     Compute $\phi(t)$
6:     $i \leftarrow 1$
7:     **while** $(i \leq n) \wedge (schedulable = true)$ **do**
8:       **if** $(t \geq r_i^1)$ **then**
9:         Compute $c_i(t)$
10:         Compute $d_i(t)$
11:         **if** $((c_i(t) > d_i(t)) \vee$
          $((t > r_i^1) \wedge (c_i(r^-(t)) > 0) \wedge (\phi(r^-(t)) \neq \tau_i) \wedge$
          $((t - r_i^1) mod T_i = 0))$ **then**
12:           $schedulable \leftarrow false$
13:         **end if**
14:       **end if**
15:       $i \leftarrow i + 1$
16:     **end while**
17:     $t \leftarrow r^+(t)$
18:     $G \leftarrow G \cup \{r^+(t)\}$
19: **end while**

IV. TIME TRIGGERED OFFLINE SCHEDULER

A. Time Triggered Approach

When dealing with schedulers at runtime, there are two approaches for triggering the tasks they manage [17]. In the usual event triggered (ET) approach the scheduler, triggered by external interruptions, selects according to an online scheduling algorithm, the next task to execute among the ready task list. In the time triggered (TT) approach the scheduler, triggered at predefined time instants, finds the next task to execute in a scheduling table built offline. Defining these time instants requires, a complete understanding of the system and of the environment it will operate in. Since we perform offline schedulability analysis that produces a scheduling table, the TT approach is the best suited for implementing our scheduler. This approach compared to ET online scheduler has the following main advantages. It prevents from exploring, online, the ready task list whose length varies according to the scheduler calls, in order to select the next task to execute. In addition, it prevents from managing online priority inversions and deadlocks since they have been taken into account during the offline schedulability analysis. Consequently, as already mentioned in the schedulability analysis, and as it will be shown afterwards at runtime, the time triggered offline scheduler is greatly simplified and deterministic compared to usual online scheduler, since its cost does not vary and is easily determined.

Among the TT schedulers, the most known are those that are called periodically. A periodic timer calls the scheduler at a predefined period that is at, best, the greatest common divisor of the task periods to prevent release and completion time misses. The main drawback of this approach is that the scheduler may be called more than necessary. The second kind of TT scheduler is called only at appropriate time instants [18], [19]. Actually, these time instants are those stored in the scheduling table.

B. Runtime Scheduler

At runtime, no task has to be selected by the scheduler since this selection has already been performed offline during the schedulability analysis. However, some actions are necessary in order to actually execute the tasks. Basically, we use the second kind of TT scheduler. The scheduling table contains in every entry the duration between two consecutive calls of the scheduler as well as the task to execute and its status. This duration is used to initialize a unique timer. This timer will interrupt, an infinite loop performing a `nop` operation, every time it reaches zero. The interruption routine is based on the algorithm 2.

This routine uses two tables as input, the scheduling table already mentioned and an additional task table that holds the context of every task. The latter table holds also the context of a specific task, called "idle". This is an infinite loop running the processor when no task runs it. Once the timer interrupt occurs, this routine is called, and immediatly loads the timer with the duration until the next call of the routine. It then updates the previous, current, and next indexes of the scheduling table, and reads the scheduling table entry at the current index. Then, using the status, a test is made on the previous executed task: if it was preempted, the routine proceeds by saving its context in the corresponding tasks table entry, otherwise (the task completes its execution, or the idle task was running) it skips the context saving. Next, it verifies if the current task was already preempted, and now must be resumed. Therefore, it retrieves its context stored in the tasks table and then restores it, otherwise it directly executes a new instance of the task. Finally, it returns from interrupt.

V. PERFORMANCE EVALUATION

A. Hardware Experimental Conditions

In order to evaluate the time triggered offline scheduler proposed in the previous sections, we sought a processor that complied with assumptions we made, i.e. neither cache nor complex pipeline or specific internal architecture features. Actually, we will use processors with a usual three stages pipeline that does not affect the performances of our approach. If we take into consideration other parameters of industrial field, the ARM Cortex-M4 processor seems to be a good choice for evaluating our scheduler. In addition to its fairly cheap cost, it is a widely spread and very popular processor in the embedded world. There are many chip vendors offering microcontrollers based on the

Algorithm 2 Interruption routine $INT\_HANDLER$

1: $INPUT1$: scheduling table $T_{SCHED}[SIZE\_T]$
2: $INPUT2$: tasks table $T_{TASKS}[TASKS\_COUNT]$
3: /*load the timer $TIMER$ with time duration stored in $T_{SCHED}[i].DURATION$ and start counting*/
4: $LOAD\_TIMER(TIMER, T_{SCHED}[i].DURATION)$
5: $START\_COUNT(TIMER)$
6: /*update scheduling table indexes i, i_prev, and i_next*/
7: $i\_prev \leftarrow i$
8: $i \leftarrow i\_next$
9: **if** $i = SIZE\_T - 1$ **then**
10:    $i\_next \leftarrow I\_PERM$
11: **else**
12:    $i\_next \leftarrow i + 1$
13: **end if**
14: /*if there is a preemption during the execution of $T_{SCHED}[i\_prev].CODE$...*/
15: **if** $(T_{TASKS}[T_{SCHED}[i\_prev].ID].OVER == FALSE) \wedge (T_{SCHED}[i].ID \neq T_{SCHED}[i\_prev].ID) \wedge (T_{SCHED}[i\_prev].ID \neq IDLE) \wedge (i \neq 0)$ **then**
16:    /*... we save the context of the preempted task $T_{SCHED}[i\_prev].CODE$ in the $CONTEXT$ field of the task table T_{TASKS} element that corresponds to the preempted task identifier*/
17:    $SAVE\_CONTEXT(T_{SCHED}[i\_prev].CODE, T_{TASKS}[T_{SCHED}[i\_prev].ID].CONTEXT)$
18: **end if**
19: /*if the execution of $T_{SCHED}[i].CODE$ was preempted...*/
20: **if** $(T_{SCHED}[i].STATUS == r) \wedge (T_{TASKS}[T_{SCHED}[i].ID].OVER == FALSE)$ **then**
21:    /*...then restore its context stored...*/
22:    $RESTORE\_CONTEXT(T_{TASKS}[T_{SCHED}[i].ID].CONTEXT)$
23:    /*... and set it to be executed. */
24:    $EXECUTE(T_{SCHED}[i].CODE)$
25: **else**
26:    /*...otherwise, the execution of $T_{SCHED}[i].CODE$ was not preempted...*/
27:    **if** $T_{SCHED}[i].STATUS == d$ **then**
28:      /*...so we execute the code direclty. */
29:      $EXECUTE(T_{SCHED}[i].CODE)$
30:    **end if**
31: **end if**
32: /* the interruption routine ends, and the task will execute*/
33: $RETURN\_FROM\_INTERRUPT()$

Cortex-M4 architecture which include the usual peripherals for communicating with the outside (USB, ethernet, etc.), efficient timers, flash storage, memory, etc.

The Cortex-M4 is also deterministic in terms of execution cycles. This is a crucial issue when we want to mesure precisely the cost of programs, not only the scheduler but also the tasks themselves. Indeed, we need to mesure precisely the WCET of the tasks without any approximation, whether they are realistic ones provided by industry or synthetic ones (simple time consuming loops). The Cortex-M4 also provides an integrated debug unit, very useful for monitoring the execution of the code, and for introducing breakpoints between the start and end of a program to measure its duration. There are two modes of execution in this processor, a privileged one with a seperated stack pointer and set of registers that we use for the scheduler context, and an unprivileged one that we use for the tasks.

For our evaluation, we chose the LPC4088 microcontroller based on the Cortex-M4. It is proposed by the NXP company and available inside a development board proposed by the Embedded Artists company. The LPC4088 provides a set of four hardware timers and a clock configuration unit to set their running periods. We use a common clock for the CPU and timers in order to avoid drifts between them. We use only two timers, one for the time triggered behaviour, configured in high priority interrupt mode and one for time measurement purposes only used for reading the elapsed time and reset afterwards. Notice that timers and in general peripherals use a dedicated peripheral bus which helps to reduce access delays when setting the timers.

B. Experiments

In order to illustrate the benefit of accounting for preemption and scheduler costs, we create two different task sets, and use an usual offline schedulability analysis based on the RM algorithm that does not account for preemption and scheduler costs, to generate their respective scheduling tables. These scheduling tables are used in a time triggered offline scheduler to run the task sets on the Cortex-M4 processor of the LPC4088 microcontroller. This scheduler written in C and assembly implements the algorithm 2. Then, we consider three different deadline miss scenarios. These experiments show the limitation of this schedulability analysis that neglects preemption and scheduler costs. Finally, we propose for each scenario, an offline schedulability analysis improved by taking into account the preemption and scheduler costs. These costs were measured on the Cortex-M4 processor. We obtained 260 cycles for the scheduler, 28 cycles for storing the context, and 28 cycles for restoring the context. These values corresponds respectively to $21\mu s$, $2.3\mu s$ and $2.3\mu s$ with a clock of $12Mhz$ that we set with the LPC4088 microcontroller clock configuration unit. When the task set is running, a logging code allows the measurements of release, completion, preemption and resume times for

each task. These measures are used to display the runtime timing diagram describing the runtime schedule of the tasks. They are depicted in figures presented afterwards, such that the first row shows preemption and scheduler costs in black along with the idle task in grey, while the following rows show the scheduling of the tasks ordered from the highest to the lowest priority. In order to easily modify its WCET, each task is built as a loop, with different stop condition values. Moreover, the duration of the scheduler is measured and can be increased by a waiting loop, to possibly change the proportion between the scheduler duration and task durations.

For every instance of a task τ_i, $CE_i = C_i - CP_i$ denotes its effective WCET. As described in section III-B, CP_i includes the cost for storing the context when a task is released while preempting another task, the cost of the scheduler, and the sum of all the cost α due to the preemptions occuring in this instance. If $CP_i = X\%C_i$ then $CE_i = (100 - X)\%C_i$. Notice that the CE_i value is a time upperbound shorter than the one obtained during the WCET analysis, and used during our offline schedulability analysis. $U = \sum_{i=1}^{n} \frac{C_i}{T_i}$ denotes the utilization factor of a set of n tasks. $UE = \sum_{i=1}^{n} \frac{CE_i}{T_i}$ denotes the runtime utilization factor of the same set of n tasks, when they are running at their effective WCETs CE_i.

1) First task set: the task set given in the table I contains three tasks t_1, t_2 and t_3. It was scheduled during the offline schedulability analysis with an RM algorithm. The corresponding timing diagram is given in figure 3. Because of its greatest period, t_3 is assigned the lowest priority by the RM algorithm and therefore is preempted five times during each instance by the two higher priority tasks t_1 and t_2. t_2 is the middle priority task and is preempted only one time for each instance. Consequently t_3 and t_2 must pay respectively five restore costs and one restore cost. The highest priority t_1 is not preempted but must pay the cost for storing the context of preempted tasks whether t_2 or t_3, and the cost of the scheduler at each release time. For this task set $U = 0.9833$. This utilization factor has been intentionally choosen close to one to illustrate the second scenario given afterwards, leading to a deadline miss.

Table I
FIRST TASK SET

Tasks	r_i^1	C_i	D_i	T_i
t_1	30	20	50	50
t_2	20	25	100	100
t_3	0	100	300	300

• *First scenario high priority task deadline miss:* here we show that even a high priority task t_1 that is not preempted


Figure 3. Task set 1 offline timing diagram

can miss its deadline if we do not precisely account for the cost of context storing of the tasks it preempts as well as the cost of the scheduler. We assume that t_2 and t_3 are running at CE_2 and CE_3, each 50% of their respectives C_i. Also, we assume that the cost $CP_1 = 5\%C_1$. In this case the runtime utilization factor is $UE = 0.67$. Therefore, if the high priority t_1 exceeds 95% of its C_1 assigned during the offline analysis, it misses its deadline D_1. Note that the lower priority tasks t_2 and t_3 do not miss their respective deadlines. This scenario was recorded with the logging tool which produces the runtime measured diagram shown in figure 4.
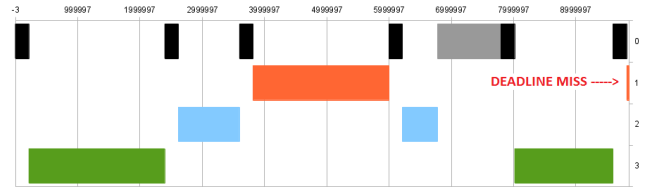

Figure 4. Scenario 1 runtime measured timing diagram

As shown on the offline timing diagram given in figure 5 where the preemption and scheduler costs are taken into account, t_1 does not miss its deadline, and a new instance of t_1 is released at time 80, whereas it misses its deadline in the runtime measured timing diagram shown in figure 4.


Figure 5. Scenario 1 offline timing diagram with scheduler and preemption costs

• *Second scenario low priority task deadline miss:* here we zoom on the last part of figure 3, shown in figure 6. When t_3 completes its instance, there is still available time before t_1 is released. However, since t_3 is preempted five

times, it pays five preemption and one scheduler costs. When these costs are neglected and are sufficiently large a deadline miss occurs as it is depicted in the runtime measured timing diagram 7.
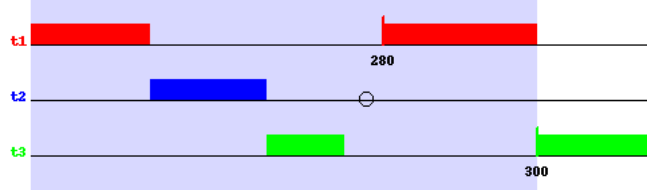


Figure 6. Task set 1 offline timing diagram

Figure 7. Scenario 2 runtime measured timing diagram

As shown on the zoomed offline timing diagram given in figure 8 where the preemption and scheduler costs are taken into account, t_3 misses its deadline, and the utilization factor is greater than one. Consequently, the task set is a priori not schedulable.

Figure 8. Scenario 2 zoomed offline timing diagram with scheduler and preemption costs

The complete offline timing diagram with preemption and scheduler costs is shown in figure 9.

2) Second task set: the task set described in the table II and in the data dependence graph shown in figure 10, contains three dependent tasks t_1, t_2 and t_3 and one independent task t_4. It was scheduled during the offline analysis with an RM algorithm. The corresponding timing diagram is given in figure 11. Because of its greatest period, t_4 is the lowest


Figure 9. Scenario 2 offline timing diagram with scheduler and preemption costs

Table II
SECOND TASK SET

Tasks	r_i^1	C_i	D_i	T_i
t_1	30	50	250	250
t_2	120	75	250	250
t_3	200	20	250	250
t_4	0	500	3000	3000

priority task. The release dates of higher priority tasks are offseted to avoid mutual preemption, but they have to pay the cost for storing the context of the only preempted t_4 and the cost of the scheduler at each release date. Due to this configuration (higher priority tasks are more frequently released than the low priority task t_4) several preemptions occurs for one instance of t_4. Such situation occurs when for instance a background task is frequently preempeted by sensor, actuator, and control tasks. The utilization factor of this task set $U = 0.7467$ is lower than the utilization factor of the first task set.


Figure 10. Task set 2 data dependence graph

Figure 11. Task set 2 offline timing diagram

• *Third scenario where preemptions produce other preemptions:* in this scenario we show that when several preemptions occur during one instance of the low priority task t_4, these preemptions costs can sufficiently delay the

task and cause additional preemptions in a cascading effect, resulting in a deadline miss. Here the task t_4 is preempted fifteen times according to the offline schedulability analysis shown in the timing diagram given in figure 11. Thus, it pays a large value CP that causes the task execution to continue beyond its offline completion time at instant 1230 pointed by the circle in the zoomed offline timing diagram given in figure 15, and pointed at position ① on the zoomed measured timing diagram in figure 14. Moreover, the task t_4 proceeds untill $r_1^6 = 1280$ the sixth release time of the higher priority task t_1. This causes a new preemption for t_4 shown at position ② in the zoomed measured timing diagram 14, and because there is no scheduling table entry afterwards to resume from this additional preemption, as shown after position ③ in the zoomed measured timing diagram 14, t_4 misses its deadline. In this case the runtime utilization factor $UE = 0.7733$ is not very close to one. However, a deadline miss occurs because of a lack in the scheduling table of the t_4 resume entry, corresponding to this additional preemption.


Figure 12.   Task set 2 zoomed offline timing diagram


Figure 13.   Scenario 3 runtime measured timing diagram


Figure 14.   Scenario 3 zoomed runtime measured timing diagram

As shown on the offline timing diagram given in figure 15 where the preemption and scheduler costs are taken into account, t_4 does not miss its deadline, whereas it misses its deadline in the runtime measured timing diagram shown in figure 14.


Figure 15.   Task set 2 offline timing diagram with scheduler and preemption costs

With the proposed approach, we avoid three different deadline miss scenarios by introducing in the schedulability analysis preemptions and scheduler costs. We use a methodology composed of three steps. In the first step we compute the scheduling with an usual offline schedulability analysis that neglects the preemption and scheduler costs. Then in the second step we compare these results with the scheduling measured by running the task set on the Cortex-M4 processor, and we observe deadline misses. Finally, we show that the improved offline schedulability analysis taking into account preemption and scheduler costs, avoid deadline misses of scenario one and three, and predicts the deadline miss of scenario two.

VI. CONCLUSION AND FUTURE WORK

We proposed a schedulability analysis for data dependent periodic tasks which precisely accounts for preemption and scheduler costs. The scheduling table produced by this analysis is exploited in a time triggered offline scheduler which is perfectly suited for time critical embedded systems, since it guarantees that no deadline misses occur in accordance with the schedulability analysis. We evaluated this scheduler on an ARM Cortex-M4 bare metal processor, and showed that it is able to schedule correctly set of tasks that miss their deadline when preemption and scheduler costs are neglected.

As future work we plan to extend the present approach to real-time multiprocessor scheduling. We plan also to take into account more complex processor architectures including caches which involve CRPD (cache-related preemption delay) costs that are more complex to master.

REFERENCES

[1] The simulink tool page: http://www.mathworks.com/products/simulink/.

[2] The scade tool page: http//www.esterel-technologies.com/products/scade-suite/.

[3] C.L. Liu and J.W. Layland. Scheduling algorithms for multiprogramming in a hard real-time environment. *Journal of the ACM*, 20(1):46–61, January 1973.

[4] A. Burns, K. Tindell, and A. Wellings. Effective analysis for engineering real-time fixed priority schedulers. *IEEE Trans. Softw. Eng.*, 21:475–480, May 1995.

[5] J. Echague, I. Ripoll, and A. Crespo. Hard real-time preemptively scheduling with high context switch cost. In *Proceedings of 7th Euromicro workshop on Real-Time Systems*, Los Alamitos, CA, USA, 1995. IEEE Computer Society.

[6] Giorgio C Buttazzo, Marko Bertogna, and Gang Yao. Limited preemptive scheduling for real-time systems. a survey. *Industrial Informatics, IEEE Transactions on*, 9(1):3–15, 2013.

[7] P. Meumeu Yomsi and Y. Sorel. Extending rate monotonic analysis with exact cost of preemptions for hard real-time systems. In *Proceedings of 19th Euromicro Conference on Real-Time Systems, ECRTS'07*, Pisa, Italy, July 2007.

[8] H. Chetto, M. Silly, and T. Bouchentouf. Dynamic scheduling of real-time tasks under precedence constraints. *Real-Time System.*, 2(3):181–194, sepember 1990.

[9] J. Forget, E. Grolleau, C. Pagetti, and P. Richard. Dynamic Priority Scheduling of Periodic Tasks with Extended Precedences. In *IEEE 16th Conference on Emerging Technologies Factory Automation (ETFA)*, Toulouse, France, September 2011.

[10] J. Forget, F. Boniol, E. Grolleau, D. Lesens, and C. Pagetti. Scheduling dependent periodic tasks without synchronization mechanisms. In *Proceedings of the 2010 16th IEEE Real-Time and Embedded Technology and Applications Symposium*, pages 301–310, Washington, DC, USA, 2010.

[11] P. Richard, Cottet, and M. Richard. On-line scheduling of real-time distributed computers with complex communication constraints. In *ICECCS*, pages 26–34, 2001.

[12] O. Kermia and Y. Sorel. A rapid heuristic for scheduling non-preemptive dependent periodic tasks onto multiprocessor. In *Proceedings of ISCA 20th International Conference on Parallel and Distributed Computing Systems, PDCS'07*, Las Vegas, Nevada, USA, September 2007.

[13] L. Sha, R. Rajkumar, and J. P. Lehoczky. Priority inheritance protocols: An approach to real-time synchronization. *IEEE Transactions on Computers*, 39:1175–1185, 1990.

[14] Theodore P. Baker. Stack-based scheduling of realtime processes. *Real-Time Systems*, 3(1):67–99, 1991.

[15] F. Ndoye and Y. Sorel. Monoprocessor real-time scheduling of data dependent tasks with exact preemption cost for embedded systems. In *Proceedings of 16th IEEE International Conference on Computational Science and Engieering, ICCSE'13*, Sydney, Australia, December 2013.

[16] A. Choquet-Geniet and E. Grolleau. Minimal schedulability interval for real-time systems of periodic tasks with offsets. *Theor. Comput. Sci.*, 310(1-3):117–134, 2004.

[17] H. Kopetz. The time-triggered model of computation. In *Proceedings of the IEEE Real Time Systems Symposium, RTSS'98*, pages 168–177. IEEE Computer Society, 1998.

[18] A. H. Kodancha. *Time Management in Partitioned Systems*. PhD thesis, Indian Institute of Science, Bangalore, Department of Computer Science and Automation, 2007.

[19] J. M. Pont. *Patterns for Time-triggered Embedded Systems: Building Reliable Applications with the 8051 Family of Microcontrollers*. ACM Press/Addison-Wesley Publishing Co., New York, NY, USA, 2001.